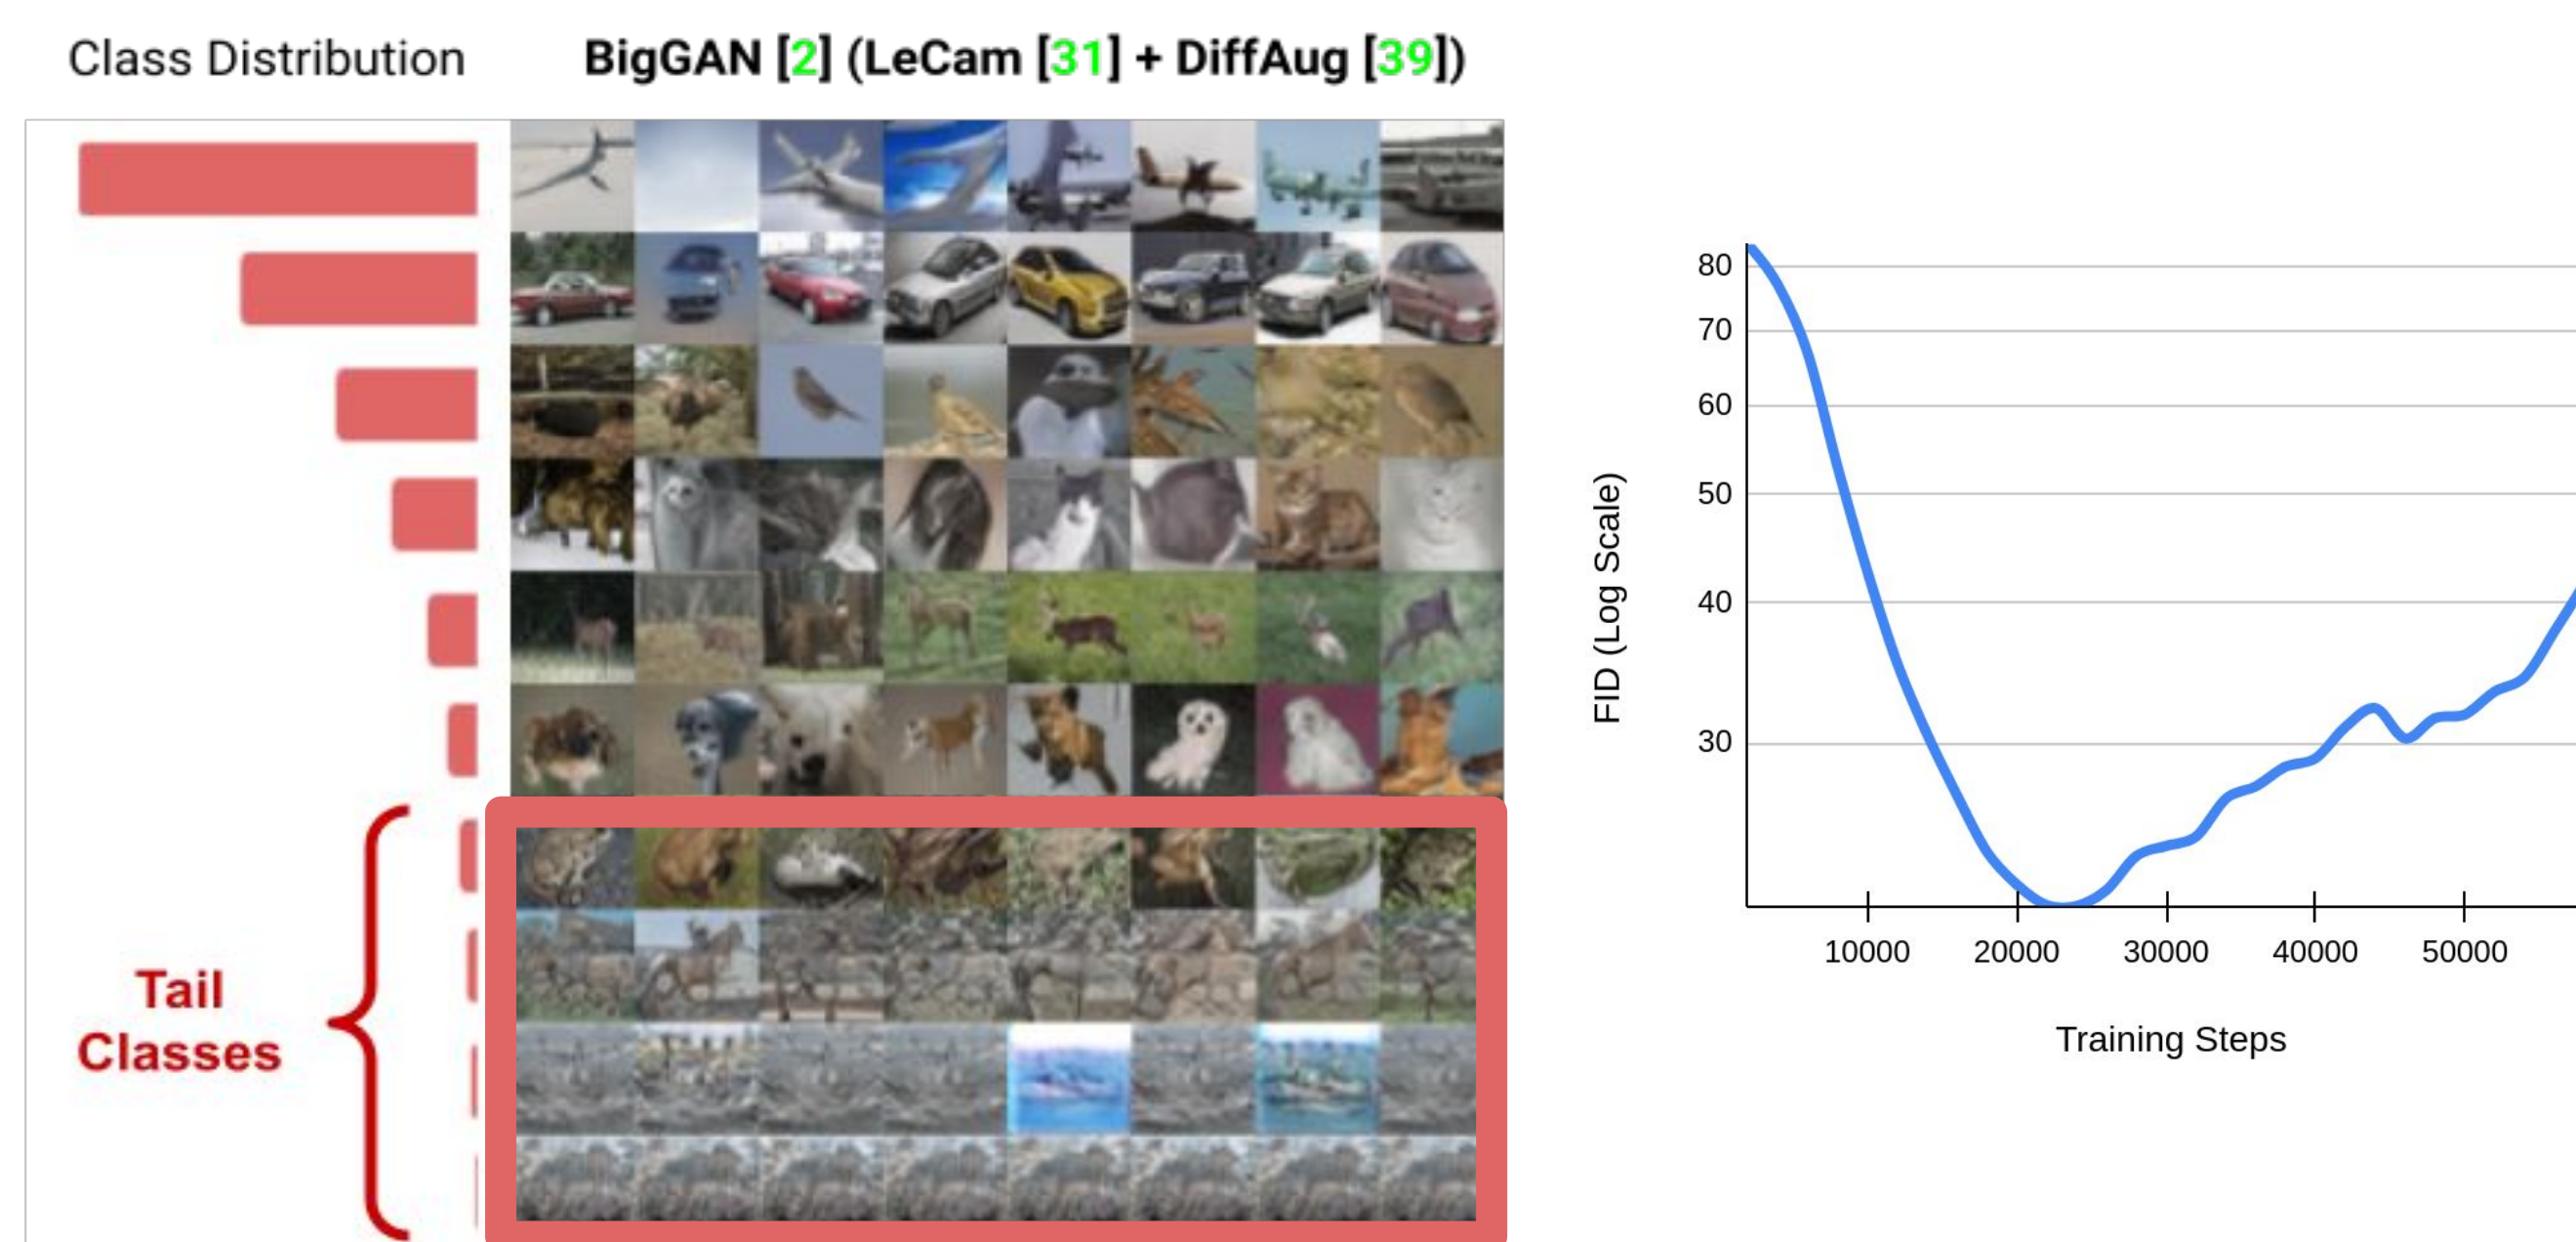
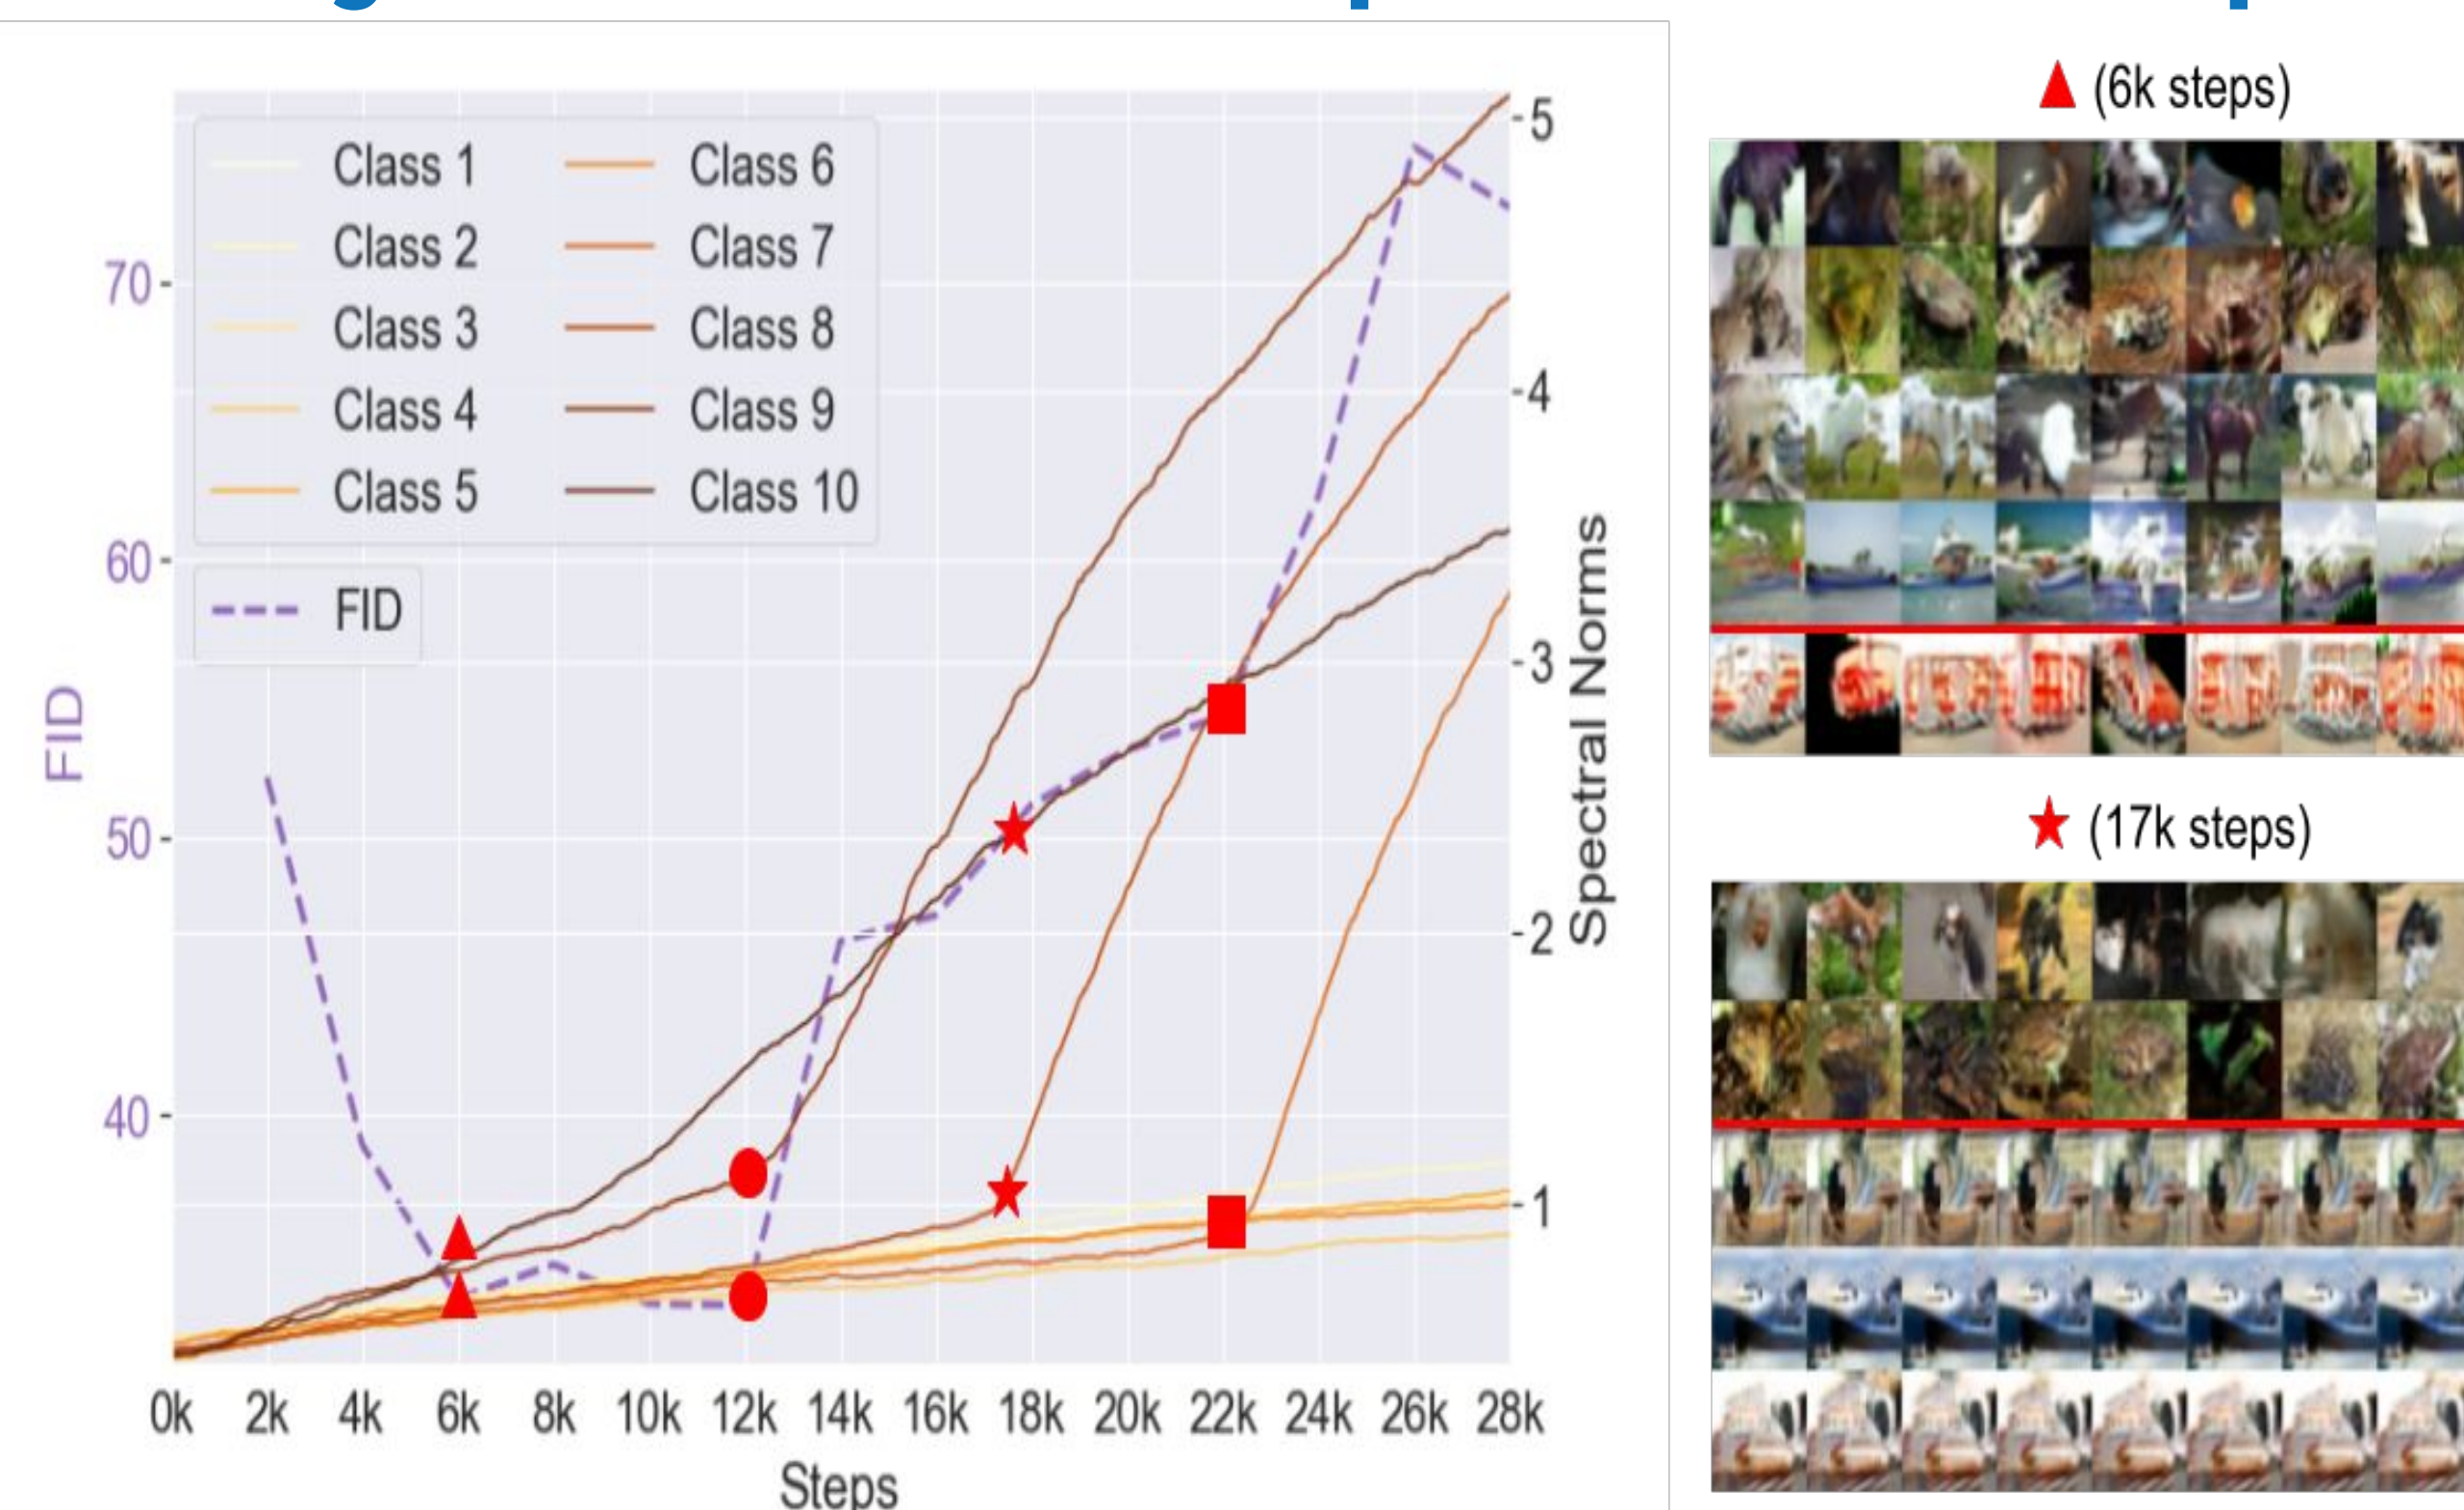


Motivation



- Long-Tailed Distribution is the **natural distribution** of categories across **species, actions** etc.
- In this work we aim to train **state-of-the-art conditional GANs** on long-tailed datasets.
- We find that the current SotA cGAN training methods suffer from **mode-collapse** while generating samples from **tail classes**.

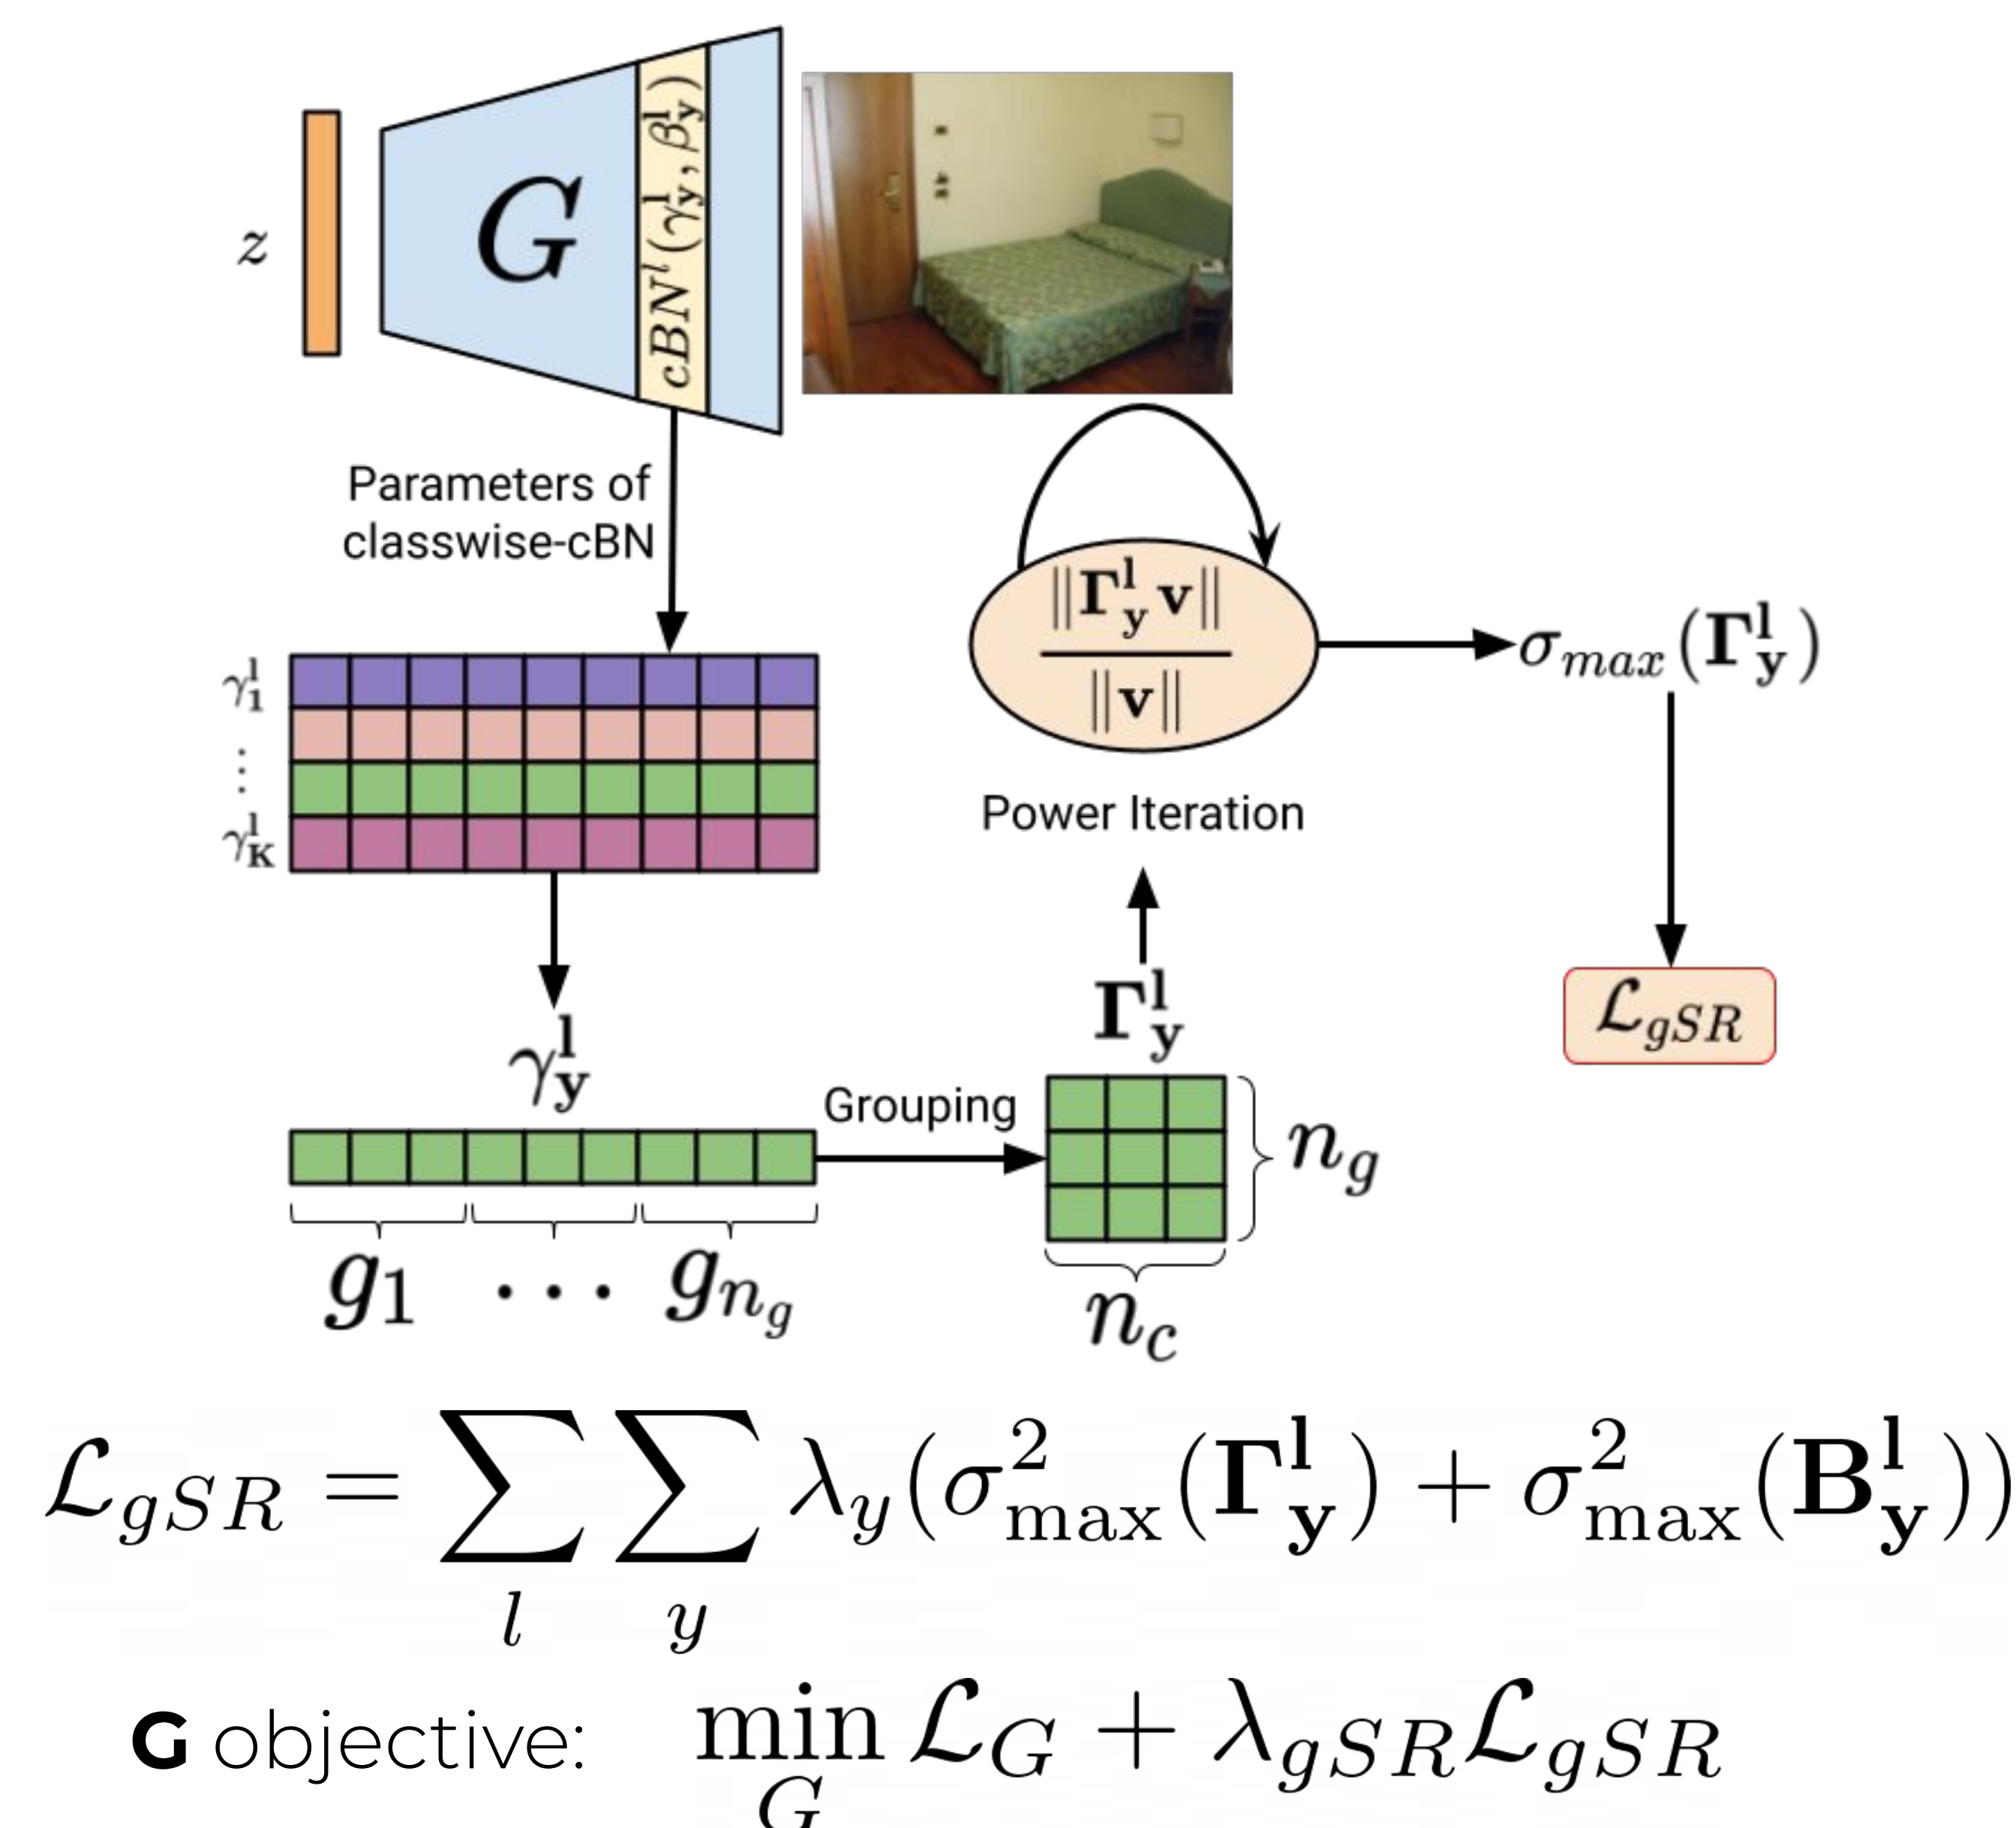
Analysis of Class-Specific Collapse



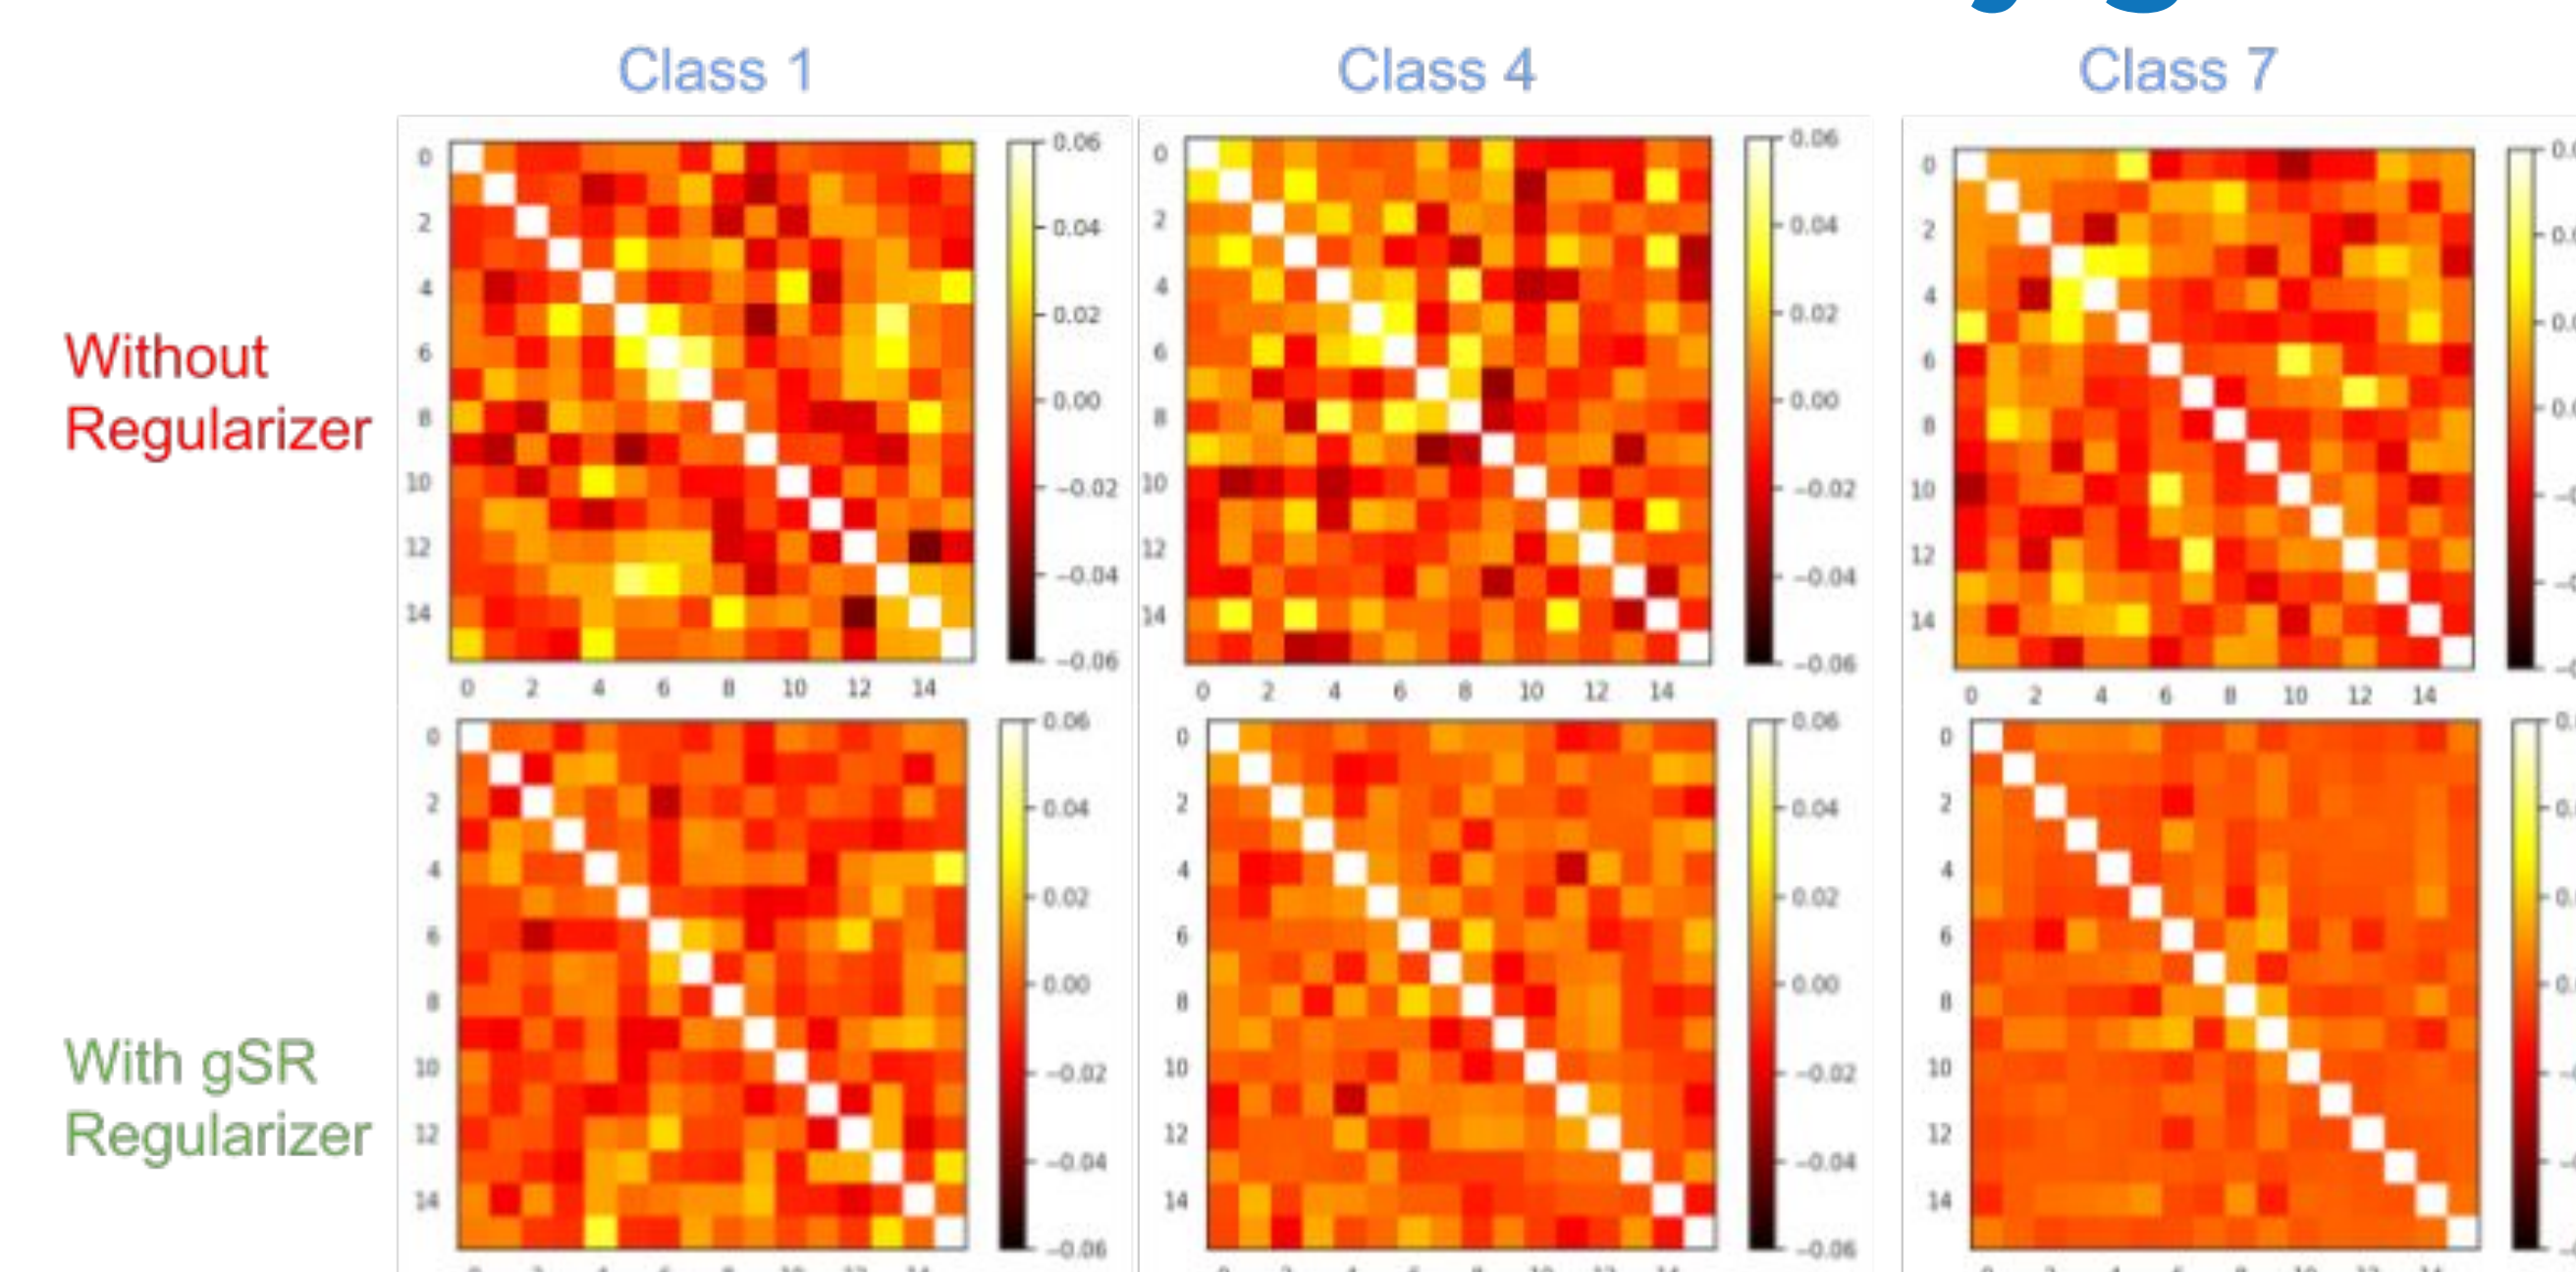
We extract the class-specific parameters from generator (G), group them and find spectral norm.

Spectral Collapse → **Class Specific Collapse**

group Spectral Regularizer (gSR)



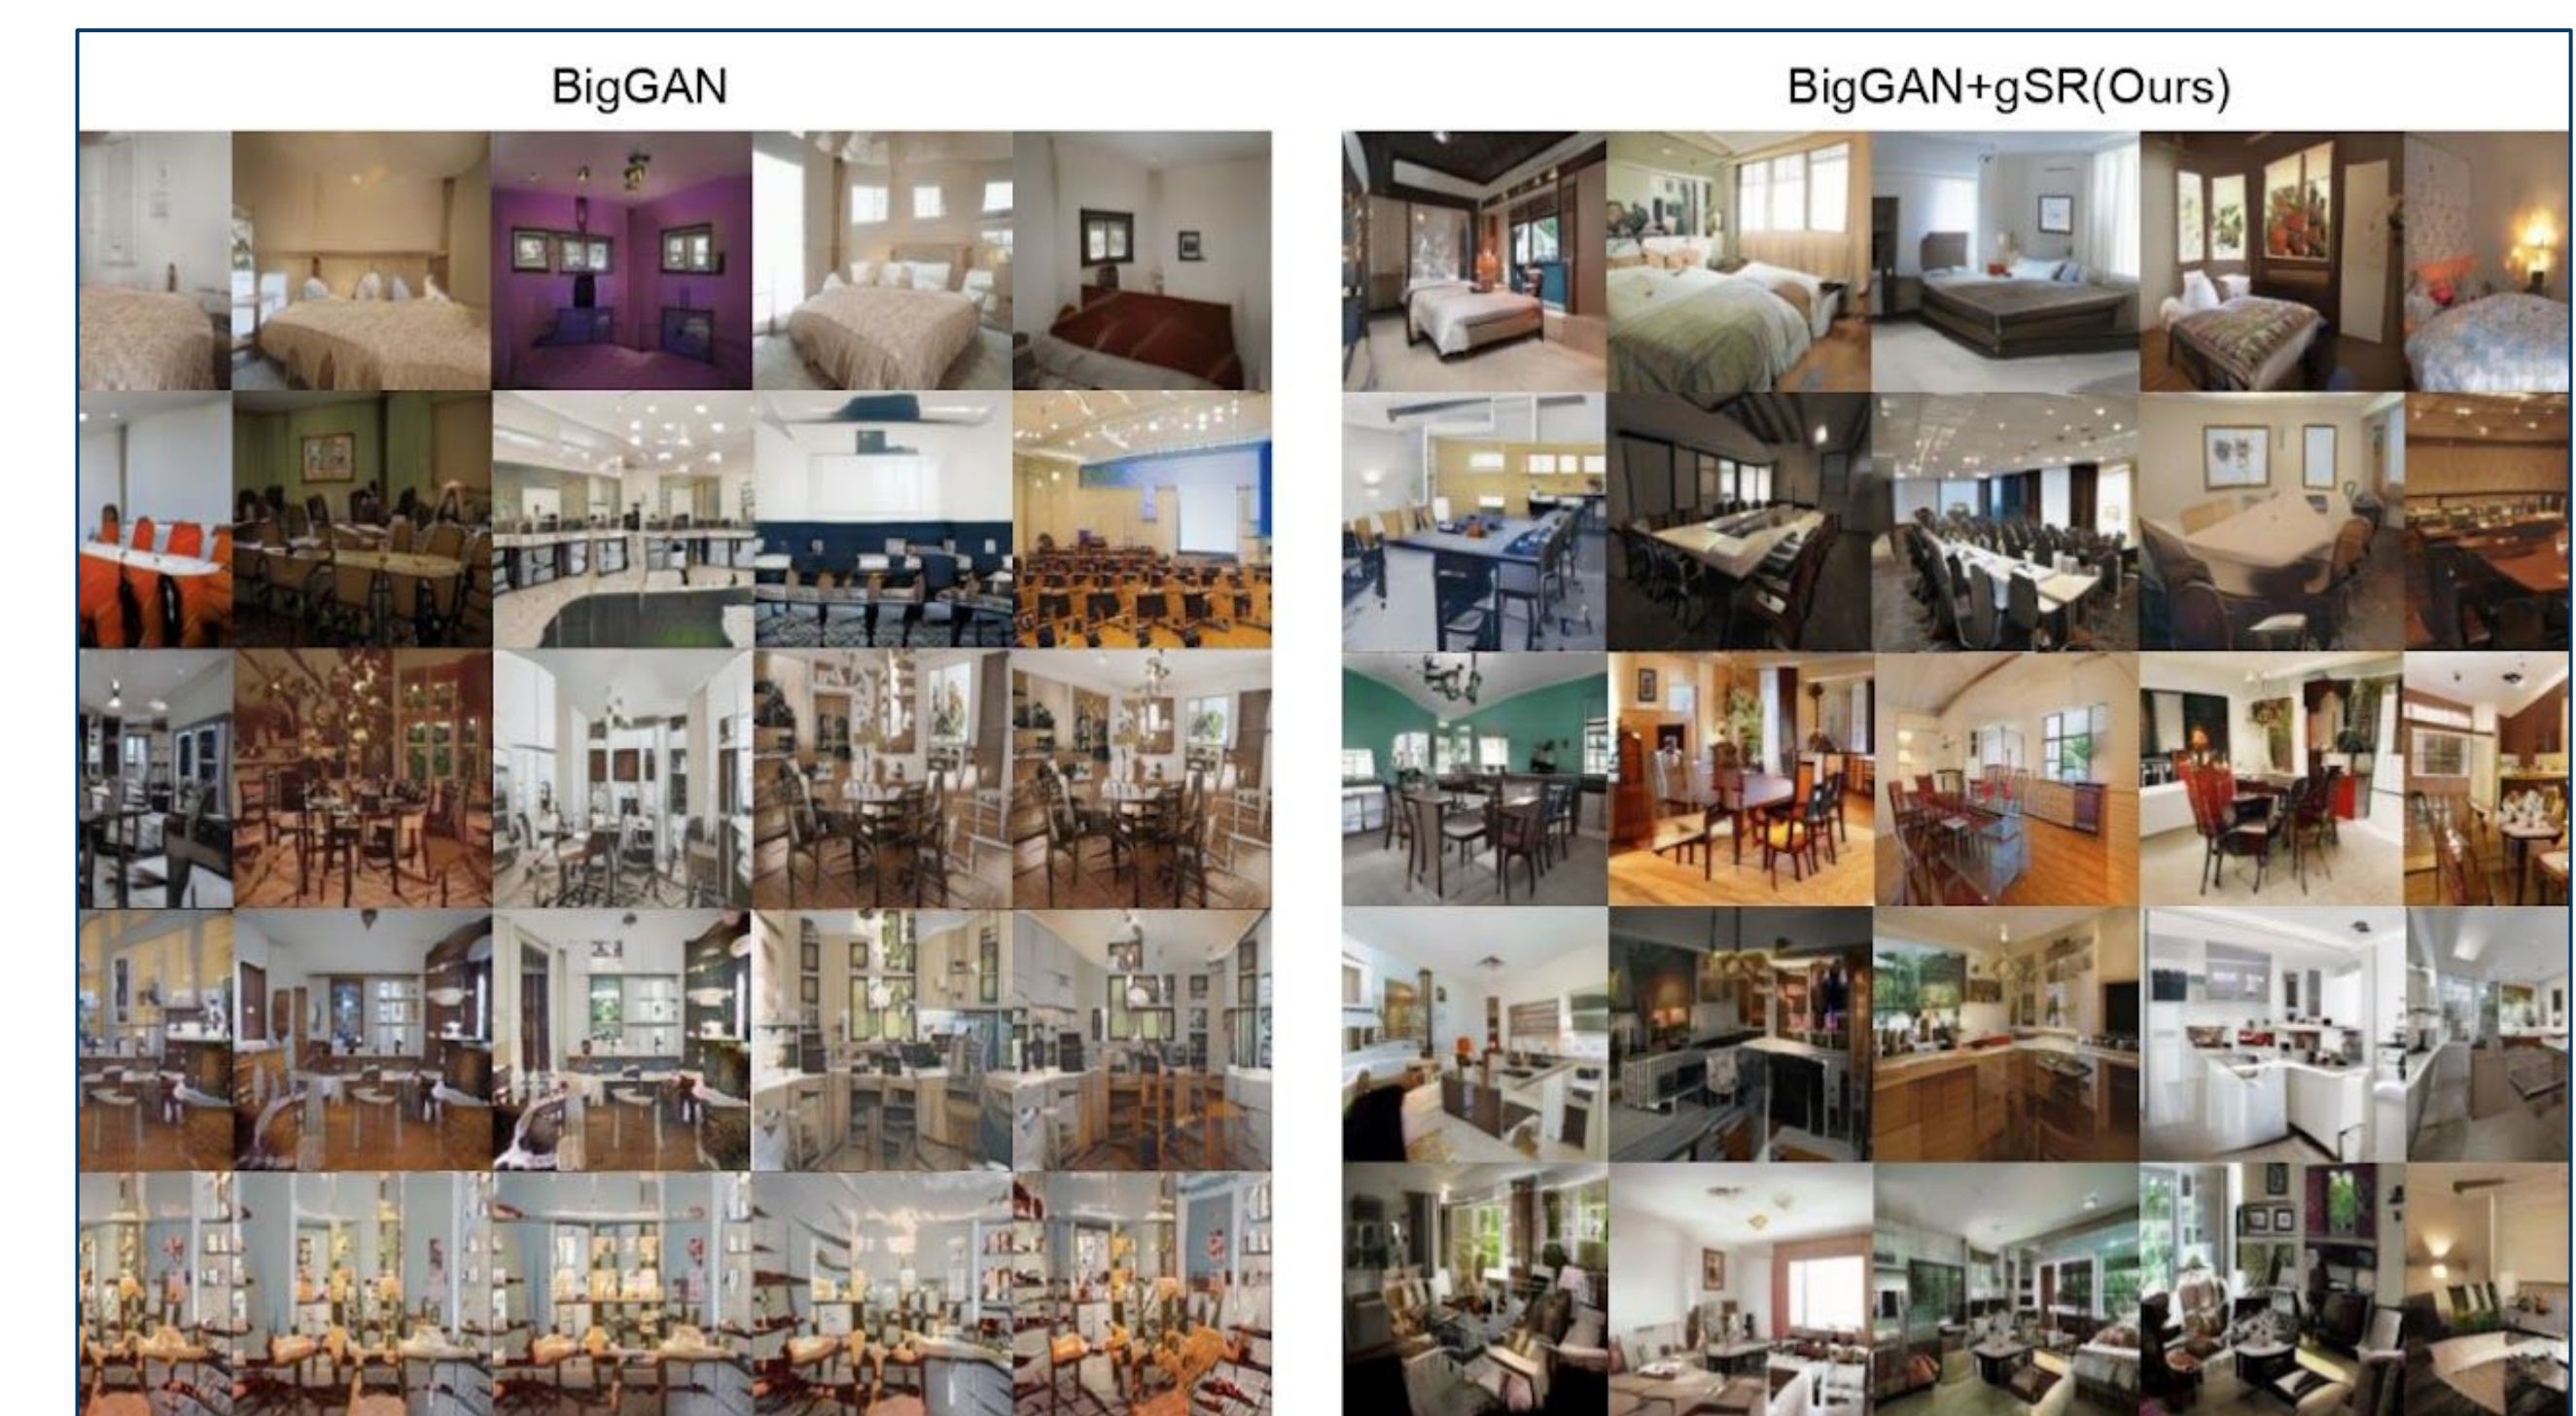

De-correlation caused by gSR



Why group Spectral Regularization?

- Addition of gSR **consistently improves** performance on imbalanced datasets.
- gSR is **inexpensive to compute** and **compatible** with BigGAN, SNGAN, StyleGAN2 architectures.

Empirical Results




Analysis of gSR (CIFAR-10 LT)

LeCam+ DiffAug	gSR	FID(↓)	IS(↑)
×	×	31.73±0.08	7.18±0.02
×	✓	27.85±0.05	7.09±0.02
✓	×	30.62±0.07	6.80±0.02
✓	✓	18.58±0.10	7.80±0.09

gSR **compliments** other regularization techniques

gSR **stabilizes** GAN training on imbalanced data, leading to improved FID.



Acknowledgement

This work was supported in part by SERB-STAR Project and Google Research Award. Harsh Rangwani is supported by PMRF fellowship.

Project: <https://sites.google.com/view/gsr-eccv22>